

lasso-ridge-elastic-regression

April 13, 2025

```
[4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
[7]: df=sns.load_dataset('mpg')
df
```

```
[7]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0    18.0          8         307.0         130.0   3504          12.0
1    15.0          8         350.0         165.0   3693          11.5
2    18.0          8         318.0         150.0   3436          11.0
3    16.0          8         304.0         150.0   3433          12.0
4    17.0          8         302.0         140.0   3449          10.5
..    ...          ...          ...          ...    ...          ...
393  27.0          4         140.0          86.0   2790          15.6
394  44.0          4          97.0          52.0   2130          24.6
395  32.0          4         135.0          84.0   2295          11.6
396  28.0          4         120.0          79.0   2625          18.6
397  31.0          4         119.0          82.0   2720          19.4

      model_year  origin  name
0             70    usa  chevrolet chevelle malibu
1             70    usa      buick skylark 320
2             70    usa    plymouth satellite
3             70    usa      amc rebel sst
4             70    usa      ford torino
..            ...    ...          ...
393           82    usa    ford mustang gl
394           82  europe      vw pickup
395           82    usa    dodge rampage
396           82    usa    ford ranger
397           82    usa    chevy s-10
```

[398 rows x 9 columns]

```
[ ]: df.drop("name",axis=1,inplace=True)
```

```
[10]: df
```

```
[10]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0    18.0          8         307.0         130.0   3504          12.0
1    15.0          8         350.0         165.0   3693          11.5
2    18.0          8         318.0         150.0   3436          11.0
3    16.0          8         304.0         150.0   3433          12.0
4    17.0          8         302.0         140.0   3449          10.5
..    ...          ...          ...          ...    ...          ...
393  27.0          4         140.0          86.0   2790          15.6
394  44.0          4          97.0          52.0   2130          24.6
395  32.0          4         135.0          84.0   2295          11.6
396  28.0          4         120.0          79.0   2625          18.6
397  31.0          4         119.0          82.0   2720          19.4

      model_year  origin
0              70     usa
1              70     usa
2              70     usa
3              70     usa
4              70     usa
..            ...     ...
393            82     usa
394            82  europe
395            82     usa
396            82     usa
397            82     usa
```

[398 rows x 8 columns]

```
[11]: df.isnull().sum()
```

```
[11]: mpg          0
      cylinders    0
      displacement  0
      horsepower    6
      weight        0
      acceleration  0
      model_year    0
      origin        0
      dtype: int64
```

```
[14]: df["horsepower"]=df["horsepower"].fillna(df["horsepower"].median())
```

```
[16]: df.dtypes
```

```
[16]: mpg            float64
      cylinders      int64
      displacement  float64
      horsepower    float64
      weight         int64
      acceleration  float64
      model_year     int64
      origin         object
      dtype: object
```

```
[17]: df["origin"].value_counts()
```

```
[17]: origin
     usa      249
     japan    79
     europe   70
     Name: count, dtype: int64
```

```
[18]: df["origin"]=df["origin"].map({"usa":1,"japan":2,"europe":3})
```

```
[20]: df["origin"].astype(int)
      df.dtypes
```

```
[20]: mpg            float64
      cylinders      int64
      displacement  float64
      horsepower    float64
      weight         int64
      acceleration  float64
      model_year     int64
      origin         int64
      dtype: object
```

```
[21]: # Now Divide x and y
      x=df.drop("mpg",axis=1)
      y=df['mpg']
```

```
[23]: # Train Test split
      from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=1,test_size=0.3)
      print(x_train.shape,x_test.shape)
```

```
(278, 7) (120, 7)
```

```
[27]: # Simple Linear Regression Model
from sklearn.linear_model import LinearRegression
regression_model=LinearRegression()
regression_model.fit(x_train,y_train)
for i,col_name in enumerate (x_train.columns):
    print(f"The coefficient for {col_name} is : {regression_model.coef_[i]}")
```

The coefficient for cylinders is : -0.3176142302799355
 The coefficient for displacement is : 0.02623748259907893
 The coefficient for horsepower is : -0.018270764913124602
 The coefficient for weight is : -0.007487750398361897
 The coefficient for acceleration is : 0.0504067346197135
 The coefficient for model_year is : 0.8470951427061368
 The coefficient for origin is : 1.519095838797505

```
[28]: # The coefficients are relatively smaller if one independent variable changes
      ↪slightly not much difference in prediction
      # sometime they called as Smoother model

from sklearn.metrics import r2_score
y_pred_linear=regression_model.predict(x_test)
r2_linear=r2_score(y_test,y_pred_linear)
print(f"The R square of Linear Regression model is: {r2_linear}")
```

The R square of Linear Regression model is: 0.8348001123742286

```
[29]: from sklearn.linear_model import Ridge
Ridge_model=Ridge(alpha=0.1)
Ridge_model.fit(x_train,y_train)

for i,col_name in enumerate (x_train.columns):
    print(f"The coefficient for {col_name} is : {Ridge_model.coef_[i]}")
```

The coefficient for cylinders is : -0.3170032101006609
 The coefficient for displacement is : 0.026213249757982955
 The coefficient for horsepower is : -0.01826325248144886
 The coefficient for weight is : -0.0074873260502131105
 The coefficient for acceleration is : 0.05036896947442607
 The coefficient for model_year is : 0.8470062938903142
 The coefficient for origin is : 1.517452828565376

```
[30]: y_pred_ridge=Ridge_model.predict(x_test)
r2_ridge=r2_score(y_test,y_pred_ridge)
print(f"The R square of Ridge Regression model is: {r2_ridge}")
```

The R square of Ridge Regression model is: 0.8348084889168357

```
[31]: from sklearn.linear_model import Lasso
lasso_model=Lasso(alpha=0.5)
lasso_model.fit(x_train,y_train)
for i,col_name in enumerate (x_train.columns):
    print(f"The coefficient for {col_name} is : {lasso_model.coef_[i]}")
```

The coefficient for cylinders is : -0.0
The coefficient for displacement is : 0.006208198888300381
The coefficient for horsepower is : -0.011058382987169605
The coefficient for weight is : -0.00698267316802309
The coefficient for acceleration is : 0.0
The coefficient for model_year is : 0.7446549520038191
The coefficient for origin is : 0.0

```
[32]: y_pred_lasso=lasso_model.predict(x_test)
r2_lasso=r2_score(y_test,y_pred_lasso)
print(f"The R square of Lasso Regression model is: {r2_lasso}")
```

The R square of Lasso Regression model is: 0.8277934716635554

```
[33]: from sklearn.linear_model import ElasticNet
elastic_model=ElasticNet(alpha=1,l1_ratio=0.5)
elastic_model.fit(x_train,y_train)
for i,col_name in enumerate (x_train.columns):
    print(f"The coefficient for {col_name} is : {elastic_model.coef_[i]}")
```

The coefficient for cylinders is : -0.0
The coefficient for displacement is : 0.005888869953667564
The coefficient for horsepower is : -0.012403874933570128
The coefficient for weight is : -0.006934550516257631
The coefficient for acceleration is : 0.0
The coefficient for model_year is : 0.7133150744603873
The coefficient for origin is : 0.0

```
[34]: y_pred_elastic=elastic_model.predict(x_test)
r2_elastic=r2_score(y_test,y_pred_elastic)
print(f"The R square of ElasticNet Regression model is: {r2_elastic}")
```

The R square of ElasticNet Regression model is: 0.8284840073256803

```
[35]: from sklearn.linear_model import LassoCV
model=LassoCV(cv=5)
model.fit(x_train,y_train)

y_pred=model.predict(x_test)
print(f"The R square of LassoCV is: {r2_score(y_test,y_pred)}")
```

The R square of LassoCV is: 0.8082805983844751

```
[37]: from sklearn.linear_model import RidgeCV
RidgeCV_model=RidgeCV(cv=5)
RidgeCV_model.fit(x_train,y_train)

y_pred_ridgecv=RidgeCV_model.predict(x_test)
print(f"The R sqaure of RidgeCV is {r2_score(y_test,y_pred_ridgecv)}")
```

The R sqaure of RidgeCV is 0.8354145247502054